

CSC3170 Assignment2

Correspondence: Yuyang Xia

October 2024

1 Introduction

- This is the second homework for CSC3170. It weighs 10% of your final grade.
- The deadline for this homework is October 27th 11:59 pm.
- This assignment includes two parts, a writing part, and a coding part, each part weighs 50%.
- Directly using AI-generated answers or other students' answers will be considered as cheating.

2 Writing part (50%)

2.1 Number of tuples in one page (10%)

1. What is the smallest size, in bytes, of a record from the following schema in a slotted page? Assume that the record header is 4B, the boolean is 1B, and the date is 8B, disregarding word-alignment.(2%)

```
name VARCHAR
student BOOLEAN
birthday DATE
state VARCHAR
```

Figure 1: Provided schema.

2. What is the maximum number of records that can be stored in a 2KB page given the schema above? Assume the slot count, free space pointer, record pointer, and record length costs 8B each. Note that you are using

the N-ary storage model (NSM) and need to use padding to conduct word-alignment where each word is 8B long. (4%)

3. If we are required to use reordering to conduct word-alignment, what is the maximum number of records that can be stored with the same setting in part 2? (4%)

2.2 Number of I/Os for query (20%)

Consider a database with a single table **Class**(id, name, instructor, size, credits), where id is the *primary key*, and all attributes are the same fixed width. Suppose **Class** has 10000 tuples that fit into 500 pages, ignore any additional storage overhead for the table (e.g., page headers, tuple headers).

Additionally, you should make the following assumptions:

- The DBMS does not have any additional meta-data (e.g., sort order, zone maps).
- **Class** does not have any indexes (including for primary key id).
- None of **Class**'s pages are already in the buffer pool.
- Content-wise, the tuples of **Class** will make each query run the longest possible (this assumption is critical for solving part 1).
- The tuples of **Class** can be in any order (this assumption is critical for solving part 2 when you compute the minimum versus the maximum number of pages that the DBMS will potentially have to read)

1. Consider the following query:

```
SELECT MAX(credits) from Class
WHERE size > 50
```

(a) Suppose the DBMS uses the decomposition storage model (DSM) with implicit offsets. How many pages will the DBMS potentially have to read from disk to answer this query? (5%, keep in mind our assumption about the contents of **Class**!)

(b) Suppose the DBMS uses the N-ary storage model (NSM). How many pages will the DBMS potentially have to read from disk to answer this query? (5%, keep in mind our assumption about the contents of **Class**!)

2. Consider another query:

```
SELECT name, instructor, size from Class
WHERE id = 114514 or id = 23333
```

- (a) Suppose the DBMS uses the decomposition storage model (DSM) with implicit offsets. What is the minimum and maximum number of pages that the DBMS will potentially have to read from disk to answer this query? (5%)
- (b) Suppose the DBMS uses the N-ary storage model (NSM). What is the minimum and maximum number of pages that the DBMS will potentially have to read from disk to answer this query? (5%)

2.3 Cuckoo hashing (20%)

Consider the following cuckoo hashing schema:

- Both tables have a size of 4.
- The hashing function of the first table returns the fourth and third least significant bits: $h_1(x) = (x \gg 2) \& 0b11$ ($0b11$ is a binary representation and equals to 3 in decimal).
- The hashing function of the second table returns the least significant two bits: $h_2(x) = x \& 0b11$.
- When inserting, try table 1 first.
- When replacement is necessary, first select an element in the second table.
- The original entries in the table are shown in the figure below.

	Table 1	Table 2
00	16	
01		
10		
11		3

Figure 2: Initial status.

As shown on the left, the entries in each table are arranged from up to down.

- Which is first inserted to the tables, 16 or 3? (5%)
- Now we insert 8 and delete 16. Please draw the hash tables after the operations. (5%)
- After the operations of part 2, we further insert 19 and then insert 4. Please draw the hash tables after the operations. (5%)

4. After the operations of parts 2 and 3, what is the minimum number of insertions to cause an infinite loop? (5%)

3 Coding part (50%)

You are required to implement a classical page replacement algorithm, the *Optimal Page Replacement Algorithm*.

Note that this coding assignment is different from assignment 1. You are expected to use Java/Python/C++ for this problem instead of SQL.

3.1 Problem description

This is a simplified version of the page replacement question. In this problem, we assume that the user only wants to read the files in the disk (**i.e., a read-only task**) and the buffer pool is currently empty. Each time, the user will only read one page of the on-disk files. Given the page reading sequence of the user, you want to figure out the minimal number of pages that the DBMS must read from the disk. Formally speaking, given a page call sequence with length n and a buffer capable of storing m pages, you need to design the optimal way of storing pages in the buffer to minimize the number of reading pages from the disk.

3.2 Input format

The first line contains two integers n and m ($n, m \leq 100000$). The second line contains an array A with n integers, meaning that the i th called page is the page with id A_i ($1 \leq A_i \leq 10^5$).

3.3 Output format

One line contains one integer, representing the minimal number of reading pages from the disk.

3.4 One example

Input:

```
5 2
1 2 3 2 4
```

Output:

```
4
```

Explanation: except for the A_4 , other pages must be read from the disk.

3.5 Limitations

Time limit: 1s

Memory limit: 256MB

1-6 test points: $n, m \leq 1000$

7-10 test points: $n, m \leq 100000$